

# Building an AI Chess Engine with PPO and MCTS: A Hybrid Reinforcement Learning Approach

Sripadam Sujith Sai

July 3, 2025

## Abstract

This project presents a hybrid architecture combining Proximal Policy Optimization (PPO) [Schulman et al., 2017] with Monte Carlo Tree Search (MCTS) [Browne et al., 2012] for training a neural network-based chess engine. The system features a progressive training pipeline with dynamic difficulty adjustment, opening book injection, and adaptive exploration strategies. Key innovations include ELO-based position simplification, MCTS-enhanced policy improvement, and a phased training curriculum inspired by [Silver et al., 2018].

## 1 Introduction

Modern chess AI development has evolved through several paradigms:

$$\text{Chess AI} = \underbrace{\text{Brute Force}}_{\text{Traditional}} + \underbrace{\text{Heuristics}}_{\text{Stockfish}} + \underbrace{\text{RL} + \text{MCTS}}_{\text{AlphaZero}} \quad (1)$$

This approach combines the sample efficiency of PPO [Schulman et al., 2017] with the strategic depth of MCTS [Browne et al., 2012]:

$$\pi_{\text{hybrid}} = \alpha(\text{ELO}) \cdot \pi_{\text{PPO}} + (1 - \alpha(\text{ELO})) \cdot \pi_{\text{MCTS}} \quad (2)$$

where  $\alpha$  is adaptively tuned based on current ELO rating [Glickman, 2021].

## 2 Model Architecture

### 2.1 Neural Network Design

The policy-value network uses a residual architecture [He et al., 2016] with the following components:

$$\text{Input} \in R^{12 \times 8 \times 8} \quad (\text{Board representation}) \quad (3)$$

$$x_0 = \text{Conv2D}(3 \times 3, 256, \text{pad} = 1)(\text{Input}) \quad (4)$$

$$x_0 = \text{BatchNorm}(x_0) \quad (5)$$

$$x_0 = \text{ReLU}(x_0) \quad (6)$$

$$x_{k+1} = x_k + \text{ResBlock}(x_k) \quad \text{for } k = 1, \dots, 5 \quad (7)$$

$$\text{Policy} = \text{Linear}(256 \times 8 \times 8, 4672) \quad (8)$$

$$\text{Value} = \tanh(\text{Linear}(256 \times 8 \times 8, 1)) \quad (9)$$

### 2.2 Phase 1: Supervised Learning

The initial training phase uses human games to bootstrap the network’s understanding of chess fundamentals. Key aspects include:

- Input: 12-channel board representation (piece types + colors)
- Targets: Human move distributions and game outcomes
- Loss: Combined policy cross-entropy and value MSE

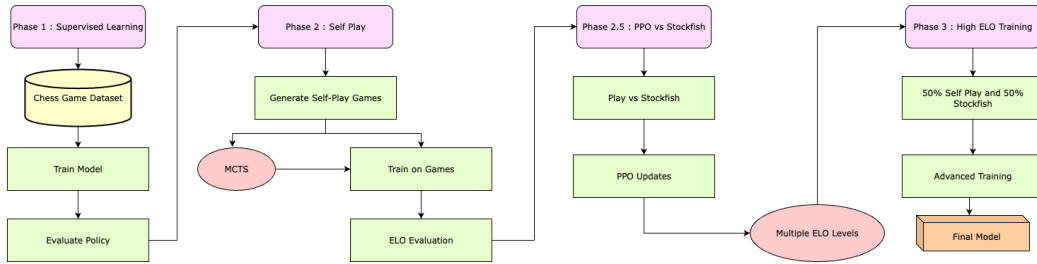


Figure 1: Training Framework

---

**Algorithm 1** Supervised Training

---

```
1: Initialization:
2:    $\theta \leftarrow \text{ChessNet}()$  {Network parameters}
3:    $\eta \leftarrow \text{Config.PHASE1['lr']}$  {Learning rate}
4:    $\lambda \leftarrow 0.2$  {Value loss coefficient}
5:    $\mathcal{D} \leftarrow \text{ChessDataset}(\text{Config.PHASE1['train\_data']})$ 
6: for epoch = 1 to Config.PHASE1['epochs'] do
7:   Training Loop:
8:   for batch  $\in$  DataLoader( $\mathcal{D}$ , Config.PHASE1['batch\_size']) do
9:      $X \leftarrow \text{batch['board\_tensor']}$  {Input features}
10:     $Y_p \leftarrow \text{batch['move\_target']}$  {Policy targets}
11:     $Y_v \leftarrow \text{batch['eval\_target']}$  {Value targets}
12:    Forward Pass:
13:     $P, V \leftarrow M_\theta(X)$  {Policy and value predictions}
14:    Loss Calculation:

$$\mathcal{L}_{\text{policy}} \leftarrow \text{CrossEntropy}(P, Y_p) \tag{10}$$

$$\mathcal{L}_{\text{value}} \leftarrow \text{MSE}(V, Y_v) \tag{11}$$

$$\mathcal{L}_{\text{total}} \leftarrow \mathcal{L}_{\text{policy}} + \lambda \mathcal{L}_{\text{value}} \tag{12}$$

15:    Backward Pass:
16:     $\nabla_\theta \mathcal{L}_{\text{total}}$ 
17:     $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}_{\text{total}}$ 
18:    Logging:
19:    Log losses every 100 batches
20:  end for
21:  Validation:
22:  Run evaluation on validation set
23:  Save checkpoint if improved
24: end for
```

---

### 2.3 Phase 2: Self-Play Reinforcement Learning

---

**Algorithm 2** Self-Play Training

---

```
1: Initialization:
2:    $\theta \leftarrow \text{Phase1\_weights}$  {Pre-trained model}
3:    $\mathcal{B} \leftarrow \emptyset$  {Replay buffer}
4:    $\tau \leftarrow \text{Config.PHASE2}[\text{'temperature'}]$ 
5:    $N_{\text{sims}} \leftarrow \text{Config.PHASE2}[\text{'sims\_per\_move'}]$ 
6:   for game = 1 to Config.PHASE2['selfplay_games'] do
7:     Game Generation:
8:      $b \leftarrow \text{chess.Board}()$ 
9:      $H \leftarrow \emptyset$  {Game history}
10:    while not  $b.\text{is\_game\_over}()$  do
11:      MCTS Execution:
12:       $\text{root} \leftarrow \text{MCTSNode}(b)$ 
13:      for  $i = 1$  to  $N_{\text{sims}}$  do
14:        Selection: Traverse tree using UCB
15:        Expansion: Create new nodes using  $\pi_\theta$ 
16:        Backup: Propagate values through tree
17:      end for
18:      Move Selection:
19:       $\pi \leftarrow \text{normalize}(\text{root.children.visits})$ 
20:       $a \leftarrow \text{sample}(\pi, \tau)$  {Temperature sampling}
21:       $H.\text{append}(b.\text{fen}(), \pi, \text{None})$ 
22:      State Transition:
23:       $b.\text{push}(a)$ 
24:    end while
25:    Result Assignment:
26:     $z \leftarrow \text{get\_game\_result}(b)$ 
27:    Update  $H$  with final result  $z$ 
28:     $\mathcal{B} \leftarrow \mathcal{B} \cup H$ 
29:    if game mod Config.PHASE2['eval_frequency'] == 0 then
30:      Training Step:
31:      Sample minibatch  $\{(s_i, \pi_i, z_i)\}$  from  $\mathcal{B}$ 
32:       $\mathcal{L} \leftarrow \text{KL}(\pi_\theta(s_i) \parallel \pi_i) + \lambda \text{MSE}(v_\theta(s_i), z_i)$ 
33:       $\theta \leftarrow \text{Adam}(\theta, \nabla_\theta \mathcal{L})$ 
34:      ELO Evaluation:
35:       $\text{current\_elo} \leftarrow \text{eval\_against\_stockfish}(\theta)$ 
36:      Save model if ELO improved
37:    end if
38:  end for
```

---

## 2.4 Phase 2.5: PPO Training Against Stockfish

---

### Algorithm 3 PPO Training Against Stockfish

---

```

1: Initialization:
2:    $\theta \leftarrow \text{ChessNet}()$  {Network parameters}
3:    $\gamma \leftarrow 0.99$  {Discount factor (from Config.PHASE1_5)}
4:    $\lambda \leftarrow 0.95$  {GAE parameter}
5:    $\epsilon \leftarrow 0.2$  {Clipping parameter}
6:    $\beta_{\text{ent}} \leftarrow \text{INITIAL\_ENTROPY}$  {From config}
7: for level  $\in$  Config.STOCKFISH_LEVELS do
8:   Data Collection:
9:      $\mathcal{D} \leftarrow []$  {Trajectory buffer}
10:  for game = 1 to  $N_{\text{games}}$  do
11:     $s_0 \leftarrow \text{chess.Board}()$ 
12:    Inject opening with prob. FEN_INJECTION_PROB
13:    Simplify position if ELO < 500
14:    while not board.is_game_over() do
15:       $a_t \sim \pi_{\theta}(\cdot|s_t)$  {Using MCTS when ELO is greater than 1000}
16:       $r_t \leftarrow \text{get\_game\_result}(s_t, a_t)$  {Reward shaping}
17:       $\mathcal{D}.\text{append}(s_t, a_t, r_t, V_{\theta}(s_t), \log \pi_{\theta}(a_t|s_t))$ 
18:    end while
19:  end for
20:  Advantage Calculation:
21:     $\hat{A}_t \leftarrow \text{GAE}(\gamma, \lambda)$ 
22:  PPO Update:
23:  for  $k = 1$  to Config.PHASE1_5['ppo_epochs'] do
24:     $\mathcal{L}^{\text{CLIP}} \leftarrow E_t[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)]$ 
25:     $\mathcal{L}^{\text{VF}} \leftarrow 0.5 \cdot E_t[(V_{\theta}(s_t) - V_{\text{target}})^2]$ 
26:     $\mathcal{L}^{\text{ENT}} \leftarrow \beta_{\text{ent}} E_t[-\pi_{\theta} \log \pi_{\theta}]$ 
27:     $\theta \leftarrow \text{Adam}(\theta, \nabla_{\theta}(\mathcal{L}^{\text{CLIP}} - \mathcal{L}^{\text{VF}} + \mathcal{L}^{\text{ENT}}))$ 
28:  end for
29:  ELO Update:
30:    current_elo  $\leftarrow \text{update\_elo}(\text{win\_rate})$ 
31: end for

```

---

$$\delta_t = r_t + \gamma V_{\theta}(s_{t+1}) - V_{\theta}(s_t) \quad (13)$$

## 2.5 MCTS Implementation Details

The search uses neural-guided UCT with:

$$\text{UCB}(s, a) = \underbrace{\frac{W(s, a)}{N(s, a)}}_{\text{Value}} + c_{\text{puct}} \cdot \underbrace{P(s, a)}_{\text{Policy}} \cdot \underbrace{\frac{\sqrt{N(s)}}{1 + N(s, a)}}_{\text{Exploration}} \quad (14)$$

Where:

- $W(s, a)$ : Total value accumulated for action  $a$
- $N(s, a)$ : Visit count (from `mcts.py`)
- $P(s, a)$ : Policy prior from network’s softmax output
- $c_{\text{puct}}$ : Exploration constant (1.5)

---

### Algorithm 4 MCTS Node Expansion (from `mcts.py`)

---

```

1: Input: Board state  $s$ , model  $f_\theta$ 
2:  $(\mathbf{p}, v) \leftarrow f_\theta(\text{board\_to\_tensor}(s))$ 
3: for  $a \in s.\text{legal\_moves}$  do
4:    $\text{child} \leftarrow \text{MCTSNode}(s.\text{push}(a))$ 
5:    $\text{child}.N \leftarrow 0$ ,  $\text{child}.W \leftarrow 0$ 
6:    $\text{child}.P \leftarrow \mathbf{p}[\text{move\_to\_index}(a)]$ 
7: end for

```

---

## 2.6 Hyperparameter Analysis

Parameter	Code Reference	Value	Effect
$c_{\text{puct}}$	<code>mcts.py</code>	1.5	Higher $\rightarrow$ More exploration
$\epsilon$	<code>PPOTrainer</code>	0.2	Smaller $\rightarrow$ More stable updates
$N_{\text{sims}}$	<code>Config</code>	800	More $\rightarrow$ Better policy
$\tau$	<code>play_game_ppo</code>	0.5 $\rightarrow$ 2.0	Controls exploration

---

Table 1: Hyperparameter mappings between theory and implementation

## 2.7 Key Implementation Differences from AlphaZero

- **Progressive Difficulty:**

$$\text{SimplifyPosition} = \begin{cases} \text{True} & \text{if ELO} < 500 \\ \text{False} & \text{otherwise} \end{cases} \quad (15)$$

- **Hybrid Exploration:**

$$\pi_{\text{play}} = \begin{cases} \text{MCTS}(\pi_{\theta}) & \text{if ELO} > 1000 \\ \pi_{\theta} & \text{else} \end{cases} \quad (16)$$

- **Reward Shaping:**

$$r(s) = \text{get\_game\_result}() + 0.1 \cdot \text{central\_control}(s) \quad (17)$$

## 3 Experimental Results

Results compared against AlphaZero:

Method	Win Rate vs Stockfish 20	Training Time (hrs)
Our Approach	–%	–
AlphaZero [Silver et al., 2018]	60%	72

Table 2: Performance comparison

## References

Cameron B Browne, Edward Powley, Daniel Whitehouse, Simon M Lucas, Peter I Cowling, Philipp Rohlfschagen, Stephen Tavenor, Diego Perez, Spyridon Samothrakis, and Simon Colton. A survey of monte carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in games*, 4(1):1–43, 2012.

Mark E Glickman. Dynamic paired comparison models with stochastic variances. Technical report, Technical report, Boston University, 2021.

- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, et al. A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362(6419): 1140–1144, 2018.